

# Multicore Multinode - A5

Marco Franzon

December, 2018

## Abstract

In this exercise, the aim is to investigate and give a measure of the latency and bandwidth between cores in the same or in different sockets or NUMA nodes of the Ulysses cluster.

## 1 Introduction - Ping Pong

Ping Pong is an Intel benchmark used to measure latency and bandwidth of a message sent between two cores. The message size keeps increasing, because initial small messages are used to estimate latency, while bigger messages are used to estimate bandwidth. We performed three tests, to measure respectively latency and bandwidth for inter-socket, intra-socket and intra-node communications between cores of Ulisse cluster.

## 2 MPI

Once loaded the module *impi-trial/5.0.1.035* and *hwloc* we ran the code using *mpirun*. Using *hwloc* we have ensured that only two cores from the same socket were used for the first run and two from two different sockets in the second. It should be noted that the results obtained for small amount of data passed are significant for the latency while the result for large data transfer are significant for the bandwidth. We loaded *impi* and *hwloc* and launched the benchmark with the following command:

```
mpirun -np 2 path-to-executable/IMB-MPI1 PingPong -iter 10000 PingPong
```

The output of the MPI benchmark is shown in Figure 1. As it can be seen the time  $t$  is constant for packets of size up to and including 16/32 bits, the average of these can be used as an estimate of the latency of the channel. On the other side, the bandwidth keeps increasing as the packet size increases, but reaches a plateau for big packet sizes. In order to have an estimate of the bandwidth, the average of the last three entries is taken. The following tables are the output of the ping pong program. First one regards two cores in the same socket. Second one is the output of two cores from different sockets.

| #bytes | #repetitions | t [usec] | Mbytes/sec |
|--------|--------------|----------|------------|
|        | 0            | 1000000  | 0.66       |
|        | 1            | 1000000  | 0.74       |
|        | 2            | 1000000  | 0.73       |
|        | 4            | 1000000  | 0.74       |
|        | 8            | 1000000  | 0.73       |
|        | 16           | 1000000  | 0.73       |
|        | 32           | 1000000  | 0.74       |
|        | 64           | 655360   | 0.78       |
|        | 128          | 327680   | 0.79       |
|        | 256          | 163840   | 0.86       |
|        | 512          | 81920    | 1.03       |
|        | 1024         | 40960    | 1.28       |
|        | 2048         | 20480    | 1.61       |
|        | 4096         | 10240    | 2.34       |
|        | 8192         | 5120     | 3.83       |
|        | 16384        | 2560     | 6.89       |
|        | 32768        | 1280     | 12.84      |
|        | 65536        | 640      | 15.57      |
|        | 131072       | 320      | 28.13      |
|        | 262144       | 160      | 50.26      |
|        | 524288       | 80       | 106.58     |
|        | 1048576      | 40       | 219.41     |
|        | 2097152      | 20       | 445.95     |
|        | 4194304      | 10       | 908.65     |

| #bytes | #repetitions | t [usec] | Mbytes/sec |
|--------|--------------|----------|------------|
|        | 0            | 903011   | 5.55       |
|        | 1            | 884632   | 5.66       |
|        | 2            | 884632   | 5.65       |
|        | 4            | 884632   | 5.65       |
|        | 8            | 884408   | 5.65       |
|        | 16           | 884408   | 5.64       |
|        | 32           | 879273   | 5.67       |
|        | 64           | 655360   | 5.69       |
|        | 128          | 327680   | 5.70       |
|        | 256          | 163840   | 5.76       |
|        | 512          | 81920    | 5.75       |
|        | 1024         | 40960    | 5.81       |
|        | 2048         | 20480    | 5.91       |
|        | 4096         | 10240    | 6.15       |
|        | 8192         | 5120     | 6.54       |
|        | 16384        | 2560     | 7.38       |
|        | 32768        | 1280     | 9.25       |

|         |     |          |        |
|---------|-----|----------|--------|
| 65536   | 640 | 901.39   | 69.34  |
| 131072  | 320 | 914.72   | 136.65 |
| 262144  | 160 | 922.85   | 270.90 |
| 524288  | 80  | 2770.42  | 180.48 |
| 1048576 | 40  | 6358.53  | 157.27 |
| 2097152 | 20  | 14007.58 | 142.78 |
| 4194304 | 10  | 28846.80 | 138.66 |

It should be noted that the results obtained for small amount of data passed are significant for the latency while the result for large data transfer are significant for the bandwidth. It's clear from the tables that the bandwidth is bigger for cores from the same socket while the latency is smaller.

### 3 Stream

Stream stand for Sustainable Memory Bandwidth Benchmark and uses four vector kernels to measure sustainable memory bandwidth (in MB/s). We set the variable OMP NUM THREADS=1 in order to perform this test on a single core, with the following results:

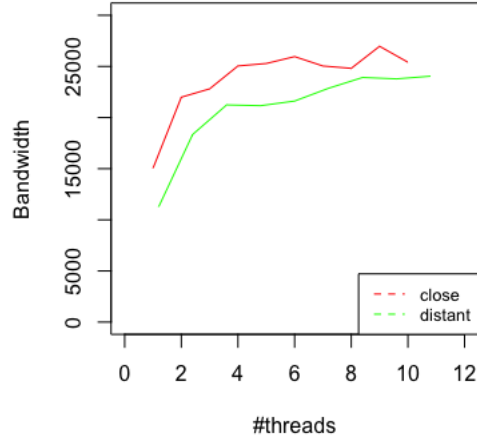


Figure 1: bandwidth for different numbers of cores, using both the memory of the same socket of the cores and the one of the other socket.

In the graph we can see that the cost on bandwidth of using a distant memory region is relevant in any case and should be avoided for better performance.

## 4 Nodeperf

In this test we run a benchmark that measure peak performance through matrix-matrix multiplication. In particular, we want to investigate the number of FLOPS attainable using the **nodeperf** benchmark tool. All the 20 cores are used in order to get an estimate of the performance. For this assignment we were supposed to compile the program `nodeperf.c` using the intel compiler but the suggested command didn't work, at first we changed the flag `-qopenmp` in `-fopenmp`, and the program compiled but while running something wrong in memory happened everytime and the cluster gave segmentation fault, in order to solve the problem we changed also the code commenting the MKL malloc line and uncommenting the line with the normal C malloc. Using 20 cores the result obtained is 448.35 GFlops, that reaching almost perfectly the peak performance.